

Supratik Padhee

+91 7978259993 | supratikpadhee.career@outlook.com | [LinkedIn](#) | [GitHub](#)

PROFILE

Data Engineer with 8+ years delivering scalable data pipelines and materialization platforms across Finance, Supply Chain, and E-commerce — currently building and operating Spark-based infrastructure processing 20TB+ daily at Intuit.

TECHNICAL SKILLS

Languages: Python, SQL, Java, Scala
Big Data & Processing: Apache Spark, Apache Hive, Delta Lake, Kafka, Avro, Apache Comet, Apache Livy, Apache Ranger
Cloud & Infrastructure: AWS S3, EMR, ECS, Lambda, DynamoDB, Glue, Kubernetes, Docker, Jenkins
Data Engineering: CDC Pipelines, Batch Materialization, Schema Registry, Data Quality, Data Governance, ETL/ELT
Frameworks & Databases: Flask, FastAPI, Elasticsearch, Neptune, MySQL, PostgreSQL, Splunk, Data Modeling

EXPERIENCE

- Intuit** May 2025 – Present
Senior Data Engineer – United Ingestion Platform Bengaluru, India
- Built and owned Spark materialisation engine on BPP — 30+ reconciler strategies, FGAC IAM, S3/Hive Metastore (HMAAS), and automated job alerting — processing 20TB+ daily across 200+ finance pipelines with high availability.
 - Replaced broadcast joins with Bloom-filter deduplication in QBO reconcilers, eliminating OOM-driven failures; demonstrated 25%+ compute cost reduction via Apache Comet with platform-wide rollout underway; and led EMR 6.x → 7.2.0 (Spark 3.5.1) upgrade across 200+ pipelines.
 - Migrated PII classification from MDR v2 → v3 async API in uip-bpp-common — converting MDR from a hard runtime dependency to a TTL-cached soft coupling, eliminating cascading failures during metadata service outages.
- Brambles** December 2021 – May 2025
Senior Data Engineer – Various Projects Bengaluru, India
- Engineered a self-service distributed computing platform on AWS EMR (Spark, Hive, Livy, Ranger, Zeppelin) — enabling data teams to self-serve pipeline execution and interactive exploration, replacing manual job request processes.
 - Built a multi-tenant onboarding API (Python/Flask on AWS ECS) enabling teams to self-provision EMR resources programmatically — reducing tenant onboarding from weeks of manual coordination to a self-serve API workflow.
- Altra Dimension Technology** September 2020 – December 2021
Cloud Data Platform Engineer Remote, India
- Designed and owned an AWS serverless identity resolution pipeline (Lambda, Glue, Neptune, DynamoDB, Kafka) — processing incremental and historical vehicle records to unify customer identities across 50+ Photo Enforcement clients in the USA and Canada, handling schema variability across jurisdictions.
- Xerox Technology Services** September 2017 – May 2020
Data Platform Engineer Bengaluru, India
- Built a distributed analytics platform on AWS EMR (Python, Flask, Spark, Hive, Livy, Ranger) for CiteWeb and Red Light Analysis — replacing manual job request processes with self-serve pipeline execution, reducing analytics turnaround from days to hours.

EDUCATION

- CDAC Electronics-city, Bengaluru** Bengaluru, India
Post Graduate Diploma in Advanced Computing (PGDAC) January 2017 – September 2017
- Silicon University, Bhubaneswar** Bhubaneswar, India
Bachelor of Technology in Computer Science and Engineering September 2011 – May 2015